

# Run Time & Performance

Kurt Schmidt

Dept. of Computer Science, Drexel University

November 15, 2016

Examples are taken from Kernighan & Pike, *The Practice of Programming*, Addison-Wesley, 1999



Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

# Intro

# Performance

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

**Objective:** To learn when and how to optimize the performance of a program.

*The first principle of optimisation is don't.*

- Knowing how a program will be used, and the environment it runs in, is there any benefit to making it faster?
- Which areas of the program should we focus on?

# Strategy

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Use the simplest, cleanest algorithms and data structures appropriate for the task
- Enable compiler options to generate the fastest possible code
  - Modern compilers optimise by default
  - Modern compilers are very good at their jobs
- Then, measure performance to see if changes are needed

# Strategy (cont.)

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Assess what changes to the program will have the most effect
  - Use a profiler to find hotspots in your code
- Make changes incrementally, re-assess
  - Consider alternative algorithms
  - Tune the code
  - Consider a lower-level language
    - Maybe just for time-sensitive components
  - Always retest your code!

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

# Profiling Code

# Profiler – `gprof`

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- A profiler watches your program run
- Reports back a bunch of information, including
  - How much time was spent in each function
  - How many times each function was called
- Use this information to find bottlenecks (areas worth improvement) in your code
- Profilers exist for most common languages, including Java, Python, and Haskell
- We will use `gprof`, which can be run with programs compiled with a gnu compiler

# Using gprof

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Use the `-p` option to compile extra information into your program:

```
$ gcc -p driver.c quicksort.c -o mySort
```

- Now run the program once, to generate metrics:

```
$ ./mySort < ins.10000 > /dev/null
```

- Note, a new data file has appeared:<sup>1</sup>

```
$ ls -ot | head -n3
total 3864
-rw-r--r-- 1 kschmidt 747 Aug 4 16:05 gmon.out
-rwxrwxr-x 1 kschmidt 9102 Aug 4 16:05 mySort
```

<sup>1</sup>It's raw data. Don't look at it yet



# Reading the gprof Report

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Supply gprof w/the name of the executable to see report on stdout:

```
$ gprof mySort
```

- Report contains 2 tables of data, with description of the information:

```
sample
=====
sample Each sample counts as 0.01 seconds.
sample
sample % cumulative self          self total
sample time  seconds seconds    calls us/call us/call name
sample 61.41    0.25    0.25      500  503.60  805.76 quicksort
sample 36.85    0.40    0.15 45721062    0.00    0.00 swap
sample  1.23    0.41    0.01
sample
sample ...
sample
=====
```

# Reading the gprof Report (cont.)

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

```
...
granularity: each sample hit covers 2 byte(s) for 2.45% of 0.41 sec

index % time    self children   called    name
-----
[1]    100.0    0.01   0.40           <spontaneous>
           0.25   0.15   500/500      main [1]
           quicksort [2]
-----
           6666956      quicksort [2]
           0.25   0.15   500/500      main [1]
[2]    98.8    0.25   0.15   500+6666956 quicksort [2]
           0.15   0.00 45721062/45721062 swap [3]
           6666956      quicksort [2]
-----
           0.15   0.00 45721062/45721062 quicksort [2]
[3]    37.0    0.15   0.00 45721062      swap [3]
-----
...

```

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

# Run-time Analysis

# Asymptotic Run Time

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- We would like to be able to compare algorithms, for arbitrarily large inputs
- We want to evaluate the growth of algorithms vs. input
  - We don't care about, e.g., processor speed, the leading coefficient, nor lower-order terms
  - E.g., linear search. If the array size doubles, so does the run-time  $\implies \Theta(n)$
  - Selection sort is a quadratic sort,  $\Theta(n^2) \implies$  doubling input size will quadruple run time (for large  $n$ )
  - Accessing an element of an array is a constant-time operation,  $\Theta(1)$ , regardless of size of array

# Lower Bound (Loose) – Little Omega

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Consider function  $T_n = 5n^2 + 17n - 12$

$$\lim_{n \rightarrow \infty} \frac{5n^2 + 17n - 12}{n} = \infty \implies T_n \in \omega(n)$$

- We say  $T_n$  is *bound below (loosely)* by  $n$
- We say  $T_n$  grows *strictly faster (asymptotically)* than  $n$
- $T_n$  *can not be bound above* by  $n$

More generally:

$$\lim_{n \rightarrow \infty} \frac{T_n}{f_n} = \infty \implies T_n \in \omega(f_n)$$

# Upper Bound (Loose)– Little Oh

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

$$\lim_{n \rightarrow \infty} \frac{5n^2 + 17n - 12}{n^3} = 0 \implies T_n \in o(n)$$

- We say  $T_n$  is *bound above (loosely)* by  $n^3$
- We say  $T_n$  grows *strictly more slowly (asymptotically)* than  $n^3$
- $T_n$  *can not be bound below* by  $n^3$

More generally:

$$\lim_{n \rightarrow \infty} \frac{T_n}{f_n} = 0 \implies T_n \in o(f_n)$$

# Asymptotically Equivalent – Theta

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

$$\lim_{n \rightarrow \infty} \frac{5n^2 + 17n - 12}{n^2} = 5 \implies T_n \in \Theta(n)$$

- We say  $T_n$  grows like  $n^2$  (asymptotically)
- $T_n$  can bound below, and above, by  $n^2$

More generally:

$$\lim_{n \rightarrow \infty} \frac{T_n}{f_n} = c, c \in \mathbb{R}^+ \implies T_n \in \Theta(f_n)$$

**Note:** This does not mean that  $T_n$  is polynomial

# Notes on Asymptotic Equivalence

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Two equivalent functions needn't look the same
- Just grow similarly
- Consider  $y = x + \sin x$ 
  - $y$  grows like a line
    - Bound above by  $y = x + 1$
    - Bound below by  $y = x - 1$
  - Clearly not a line



# Quick Observations

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

$$f \in o(g) \iff g \in \omega(f)$$

$$T \in \Theta(f) \implies T \notin o(f)$$

$$T \in \Theta(f) \implies T \notin \omega(f)$$

$$T \in \omega(f) \implies T \notin \Theta(f)$$

$$T \in o(f) \implies T \notin \Theta(f)$$

# Big-Oh, -Omega

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Upper (lower) bound, may or may not be **tight**

$$O(f) = o(f) \cup \Theta(f)$$

$$\Omega(f) = \omega(f) \cup \Theta(f)$$

We have these observations:

$$o(f) \subset O(f)$$

$$\omega(f) \subset \Omega(f)$$

Finally,

$$T \in \Theta(f) \iff T \in O(f) \wedge T \in \Omega(f)$$

# Qualitative Statements

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

$T$  is  $O(f)$

- “ $T$  grows no faster than  $f$ ”
- “ $T$  is bound above by  $f$ ”
- “ $f$  is an upper bound for  $T$ ”

$T$  is  $\Omega(f)$

- “ $T$  grows no slower than  $f$ ”
- “ $T$  is bound below by  $f$ ”
- “ $f$  is a lower bound for  $T$ ”

# Ranking of Common Functions

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

For reference, here are some common functions, in increasing order:

1 (constant)

$$n^2$$

$\log n$

$\vdots$

$\sqrt{n}$

$$n^p$$

$n$

$$c^n$$

$n \log n$

$$n^n \approx n!$$

$n\sqrt{n}$

Note,  $\log n \in o(n^p), p \in \mathbb{R}, \forall p > 0$

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

# Timing Programs

# Timing – time

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- If we can't evaluate the algorithm, we can run the program with various inputs, time each run
- Various languages may provide their own mechanism for timing from within the program
- The `time` utility takes a program, with arguments, to run
  - You now know why you type `date` to see what the time is

# Using the `time` Utility

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

```
time [options] cmd [cmd_args]
```

**options** Options to modify behavior of `time`. Must precede `cmd`

**cmd** The program run you want to time

**cmd\_args** Arguments, including options, to be passed to `cmd`

**Note:** This is a built-in in Bash, Tenex C-Shell, and others.

# time example

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

From our previous example:

```
$ time ./mySort < ins.10000 > /dev/null
```

- Output to screen is expensive
- We're not interested in that time

Output from time (to stdout:

```
real 0m7.572s
user 0m7.555s
sys  0m0.004s
```



# Description of Output

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

**real** The wall clock. **Total time elapsed**. Keeps ticking, even if your program is sliced out

**user** The actual time your program **spent running**, in **user mode**

**sys** The actual time your program **spent running**, in **kernel mode**

- The **sum of the user and sys times is probably what you want**

# time Built-in v. Utility

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- There is a utility (not a shell built-in)
  - On tux, installed as `/usr/bin/time`
  - Can also **report on other metrics**
  - Output values and format can be customised to `stderr`
- Bash, Tenex C Shell, and others have a built-in `time` command, which doesn't allow for customising the output format
  - The built-ins are a bit slicker, parsing up the command line
  - For e.g., the following produces no output (why?), but works fine w/the shell built-in

```
$ /usr/bin/time ./mySort < ins.10000 &> /dev/null
```

# Alternatives to Timing Program

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Computers have gotten fast
  - Makes it a little harder to grab good numbers
- Alternatively, we could be creative, use metrics from a profiler
  - E.g., we could count the number of calls to swap for various sized inputs to our sort
  - Or, we might use a function to compare elements in a sort, use a profiler to count the number of times items are compared
- We can use this data in the same way, plot it, see how it grows

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh

Ranking of Functions

Timing  
Programs

Program  
Growth

# Program Growth

# Estimating Program Growth

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- Run your program on various-sized inputs, collect times
  - Get a good number of points
  - Discard very small results
  - Provide inputs large enough to get past lower-order noise
- Find an upper and lower bound
- Consider  $\frac{T_n}{f_n}$  for various functions  $f$ 
  - Identify upper and lower bounds
  - Try to pinch in, get the upper and lower bounds closer to each other
  - You might not get them to meet

# Estimating Program Growth – e.g.

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Consider times  $T$  discovered for various input sizes  $n$ :

$n$	$T(n)$
10.00	3908.51
20.00	20657.40
30.00	55954.53
40.00	113992.17
50.00	198284.36
60.00	311920.28
70.00	457689.75
80.00	638156.74
90.00	855706.82
100.00	1112580.00

■  $T(n)$  appears to be increasing w/out bound

■ So,  $T$  is not constant

■ Maybe. We don't *know* this

Let's compare to  $f(n) = n$

# E.g. – Line is a Lower Bound

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

$n$	$T(n)/n$
10.00	390.85
20.00	1032.87
30.00	1865.15
40.00	2849.80
50.00	3965.69
60.00	5198.67
70.00	6538.43
80.00	7976.96
90.00	9507.85
100.00	11125.80

- $T(n)/n$  also appears to be increasing, w/out bound
- So,  $f(n) = n$  looks like a lower bound
  - I.e.,  $T(n) \in \Omega(n)$
  - If not tight, if it increases w/out bound, then  $T(n) \in \omega(n)$

Let's try  $f(n) = n^2$ :

# E.g. – Quadratic is a Lower Bound

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

$n$	$T(n)/n^2$
10.00	39.09
20.00	51.64
30.00	62.17
40.00	71.25
50.00	79.31
60.00	86.64
70.00	93.41
80.00	99.71
90.00	105.64
100.00	111.26

- $T(n)/n^2$  also appears to be increasing

- $f(n) = n^2$  looks like a lower bound

- I.e.,  $T(n) \in \Omega(n^2)$

- If not tight, if it increases w/out bound, then  $T(n) \in \omega(n^2)$

Let's try  $f(n) = n^3$ :



# E.g. – Cubic is an Upper Bound

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

$n$	$T(n)/n^3$
10.00	3.91
20.00	2.58
30.00	2.07
40.00	1.78
50.00	1.59
60.00	1.44
70.00	1.33
80.00	1.25
90.00	1.17
100.00	1.11

- $T(n)/n^3$  is decreasing
- $f(n) = n^3$  looks like an upper bound
  - I.e.,  $T(n) \in O(n^3)$
  - If not tight, if the values tend towards 0, then  $T(n) \in \omega(n^2)$
- We know that  $T$  grows no slower than a quadratic, and no faster than a cubic
- We *might* be able to improve one or both of those bounds

# E.g. $-n^2 \log n$

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Consider  $f_n = n^2 \log n$

$n$	$\frac{T(n)}{n^2 \log n}$
10.0	16.974
20.0	17.239
30.0	18.279
40.0	19.313
50.0	20.274
60.0	21.162
70.0	21.986
80.0	22.755
90.0	23.477
100.0	24.159

- $\frac{T(n)}{n^2 \log n}$  also seems to be increasing
- So, we have a new (better) lower bound
  - $T(n) \in \Omega(n^2 \log n)$

Let's try moving up a bit more:

E.g. –  $n^{2.3}$

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Consider  $f_n = n^{2.3}$

$n$	$T(n)/n^{2.3}$
10.0	19.589
20.0	21.024
30.0	22.411
40.0	23.558
50.0	24.528
60.0	25.369
70.0	26.112
80.0	26.781
90.0	27.388
100.0	27.947

- I'm comfortable saying  $n^{2.3}$  is a lower bound
  - $T(n) \in \Omega(n^{2.3})$

Let's try moving up a bit more:

# E.g. – $n^{2.5}$

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Consider  $f_n = n^{2.5}$

$n$	$T(n)/n^{2.5}$
10.0	12.360
20.0	11.548
30.0	11.351
40.0	11.265
50.0	11.217
60.0	11.186
70.0	11.164
80.0	11.148
90.0	11.136
100.0	11.126

- This looks like an upper bound

- $T_n \in O(n^2\sqrt{n})$

- Is it tight?
- Let's try something a little lower:

E.g. –  $n^{2.4}$

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

Consider  $f_n = n^{2.4}$

$n$	$T(n)/n^{2.4}$
10.0	15.560
20.0	15.581
30.0	15.949
40.0	16.290
50.0	16.587
60.0	16.845
70.0	17.074
80.0	17.279
90.0	17.464
100.0	17.633

■ Increasing, so, lower bound

■  $T_n \in \Omega(n^{2.4})$

# E.g. – Conclusion

Run Time &  
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time  
Analysis

Big-Oh  
Ranking of Functions

Timing  
Programs

Program  
Growth

- We have  $T_n$  bound below by  $n^{2.4}$  and bound above by  $n^{2.5}$
- We maybe didn't find it exactly, but we have a very good idea how this algorithm grows
- Only push in each direction while you're comfortable w/the data
- None of this is proof
  - We need to choose input size sufficiently large to get past lower-order terms
  - No way to *know*
  - But, a program, on a given computer, has a practical upper limit on input size